

SRAP: Sender-side Receiver-Aware Port selection for High-Speed Multi-Flow TCP

Shingo Hattori

shattori@hpcs.cs.tsukuba.ac.jp

University of Tsukuba

Tsukuba, Ibaraki, Japan

Osamu Tatebe (Advisor)

tatebe@cs.tsukuba.ac.jp

University of Tsukuba

Tsukuba, Ibaraki, Japan

Abstract

Achieving 400 Gbps requires aggregating multiple flows across cores rather than pushing single-flow limits. However, with 16×25 Gbps TCP flows, random ephemeral ports cause receiver-side packet steering (RSS) to concentrate flows on single CPU cores, degrading throughput from 25 Gbps to below 5 Gbps.

Unlike receiver-side solutions that suffer from cache misses and state migration overhead, SRAP transparently controls source ports at senders, ensuring collision-free mapping without runtime remapping costs or application modification.

With 16×25 Gbps flows, random assignment achieves optimal distribution with probability 1.04%, causing throughput to vary from 44.8 to 395 Gbps, while our approach consistently delivers 23.3–25 Gbps per flow with guaranteed 1:1 flow-to-core mapping.

CCS Concepts

- **Networks** → Network experimentation; Transport protocols; Network algorithms.

Keywords

traffic engineering, pacing, 4-tuple, flow collision, RSS, multi-flow

ACM Reference Format:

Shingo Hattori and Osamu Tatebe (Advisor). 2025. SRAP: Sender-side Receiver-Aware Port selection for High-Speed Multi-Flow TCP. In *Proceedings of Make sure to enter the correct conference title from your rights confirmation email (Conference acronym 'XX)*. ACM, New York, NY, USA, 2 pages. <https://doi.org/XXXXXXX.XXXXXXX>

1 Introduction

Packet pacing prevents microbursts and improves TCP performance in high-bandwidth scenarios [2, 4, 6]. Single-flow throughput is inherently limited by sequential processing constraints and CPU clock speeds. Rather than pushing against these fundamental limits, aggregating multiple flows across multiple cores provides a more scalable path to high throughput. Thus, achieving 400 Gbps requires distributing the load across multiple paced flows (e.g., 16×25 Gbps), leveraging the parallel processing capabilities of modern multi-core systems.

However, when Receive Side Scaling (RSS) maps multiple paced flows to the same CPU core, maintaining target-rate pacing becomes impossible. While RSS++ [1] attempts dynamic table adjustment at receivers, this incurs cache misses, packet reordering, and state migration overhead at 400 Gbps. This paper proposes SRAP,

a transmit-side approach that prevents flow collisions through intelligent source port selection, avoiding the runtime overhead of dynamic remapping while ensuring predictable performance.

2 Problem Analysis

RSS distributes packets across CPU cores by hashing the 4-tuple (source/destination IPs and ports) to index an indirection table mapping to receive queues [5].

Random ephemeral port assignment (typically 32768–60999) creates unpredictable flow-to-core mappings. With 16 flows and 32 cores, the probability of collision-free assignment is merely 1.04%.

When flows collide, they must share the core's processing capacity. Figure 1 shows the extreme case—all 16 flows mapping to one core, limiting each to below 5 Gbps (44.8 Gbps aggregate). With random port assignment, collision patterns vary unpredictably from ideal 1:1 mapping to such extreme concentration, making multi-flow aggregation unreliable.

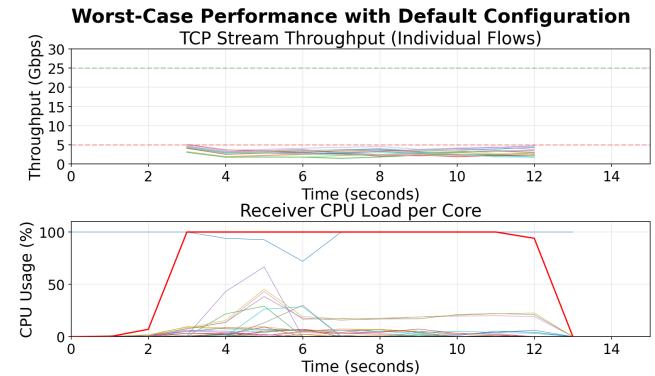


Figure 1: Random port assignment causing severe flow collisions: flows drop below 5 Gbps while one CPU reaches 100% utilization

While sender-side CPU allocation is controllable, receiver-side RSS assignment is opaque to senders. This work enables transmit-side intervention by selecting source ports to achieve predictable flow-to-core mappings.

3 Proposed Approach

SRAP controls receiver-side flow distribution by selecting source ports that hash to different CPU cores through three coordinated components (Figure 2).

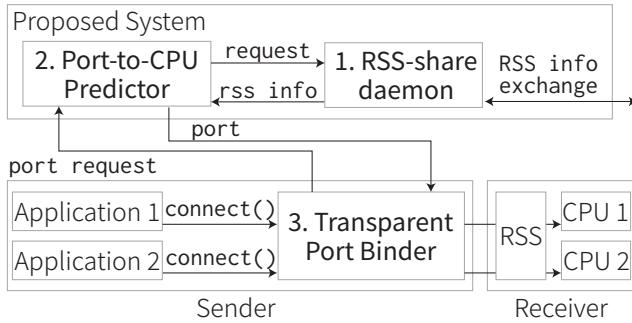


Figure 2: Proposed system architecture with three coordinated components for collision-free source port selection

(1) RSS Information Exchange: Exchanges RSS configurations (indirection tables, hash functions, IRQ mappings) with receivers for accurate port-to-core prediction.

(2) Port-to-CPU Predictor: Computes RSS hashes for potential source ports and returns collision-free ports via round-robin allocation.

(3) Transparent Port Binder: Intercepts connect() calls transparently, requests collision-free ports, and calls bind() to control the RSS 4-tuple without application modification.

By controlling source ports, SRAP prevents receiver-side flow collisions, achieving predictable multi-flow pacing without modifying applications or receivers.

4 Evaluation and Results

4.1 Experimental Setup

We evaluate between two Dell PowerEdge R6615 servers with AMD EPYC 9354P (32 cores, 3.25 GHz) and ConnectX-7 400GbE NICs running Linux 6.8.0. We use 16 parallel TCP CUBIC flows at 25 Gbps each (400 Gbps aggregate) with Linux Fair Queuing pacing [3], iperf3 with zero-copy optimizations, and 100MB TCP windows.

4.2 Evaluation

Figure 3 demonstrates SRAP's effectiveness. With our approach, all 16 flows maintain stable 23–25 Gbps throughput (top), while CPU load is evenly distributed across cores at approximately 50–65% (bottom). This contrasts sharply with the random assignment case shown in Figure 1, where flows dropped below 5 Gbps and CPUs reached 100% utilization due to collisions.

SRAP achieves perfect 1:1 flow-to-core mapping, eliminating the 99% collision probability inherent in random port assignment. Aggregate throughput consistently reaches 380–395 Gbps across all runs, compared to the 44.8–395 Gbps variance observed without control.

5 Conclusion

We demonstrated that random ephemeral port assignment causes severe RSS flow collisions, with only 1.04% probability of achieving optimal distribution for 16 flows on 32 cores, resulting in throughput varying from 44.8 to 395 Gbps.

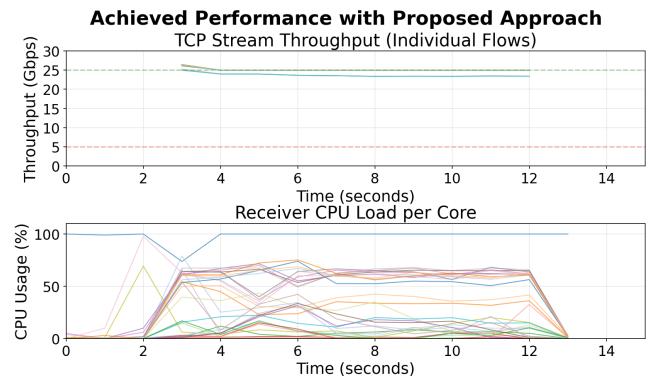


Figure 3: Achieved performance with our approach: stable 23–25 Gbps per flow with evenly distributed CPU load (50–65%)

SRAP (Sender-side Receiver-Aware Port selection) transparently controls source ports, achieving consistent 1:1 mapping and 380–395 Gbps throughput without modifying applications or receivers.

This demonstrates that source port engineering transforms unpredictable multi-flow performance into reliable high-throughput communication essential for 400G+ environments.

Acknowledgment

This work is based on results obtained from the project, “Research and Development Project of the Enhanced infrastructures for Post-5G Information and Communication Systems” (JPNP20017), commissioned by the New Energy and Industrial Technology Development Organization (NEDO), and was supported by JSPS KAKENHI Grant Number JP22H00509, JST CREST JPMJCR24R4, the Multi-disciplinary Cooperative Research Program in CCS, University of Tsukuba, and Fujitsu Limited.

References

- [1] Tom Barbette, Georgios P. Katsikas, Gerald Q. Maguire, and Dejan Kostić. 2019. RSS++: load and state-aware receive side scaling. In *Proceedings of the 15th International Conference on Emerging Networking Experiments And Technologies* (Orlando, Florida) (*CoNEXT '19*). Association for Computing Machinery, New York, NY, USA, 318–333. doi:10.1145/3359989.3365412
- [2] Neal Cardwell, Yuchung Cheng, C. Stephen Gunn, Soheil Hassas Yeganeh, and Van Jacobson. 2017. BBR: congestion-based congestion control. *Commun. ACM* 60, 2 (Jan. 2017), 58–66. doi:10.1145/3009824
- [3] Jake Edge. 2013. TSO sizing and the FQ scheduler. *LWN.net* (August 2013). <https://lwn.net/Articles/564978/>
- [4] Nathan Hanford, Brian Tierney, and Dipak Ghosal. 2015. Optimizing data transfer nodes using packet pacing. In *Proceedings of the Second Workshop on Innovating the Network for Data-Intensive Science* (Austin, Texas) (*INDIS '15*). Association for Computing Machinery, New York, NY, USA, Article 4, 8 pages. doi:10.1145/2830318.2830322
- [5] Tom Herbert and Willem de Bruijn. 2015. Scaling in the Linux Networking Stack. <https://www.kernel.org/doc/Documentation/networking/scaling.txt>. Linux Kernel Documentation.
- [6] Marcos Schwarz, Brian Tierney, Kiran Vasu, Eli Dart, Christian Esteve Rothenberg, Jeronimo Bezerra, and Italo Valcy. 2025. Recent Linux Improvements that Impact TCP Throughput: Insights from R&E Networks. In *Proceedings of the SC '24 Workshops of the International Conference on High Performance Computing, Network, Storage, and Analysis* (Atlanta, GA, USA) (*SC-W '24*). IEEE Press, 775–784. doi:10.1109/SCW63240.2024.00111

Received 20 February 2007; revised 12 March 2009; accepted 5 June 2009